

Paul Kummer

Dr. Hanku Lee

Operating Systems CSIS 430-01

10/07/2021

## Lab 05: XV6 Copy & CPS

### **Purpose of Lab**

The purpose of this lab is to show the students the difference between system calls and a regular command. The System calls require editing all the system files to update them with a new capability, whereas a command is simple and only requiring the file of the command and editing the makefile. This lab demonstrates specifically how to make a copy command and a current program status system call.

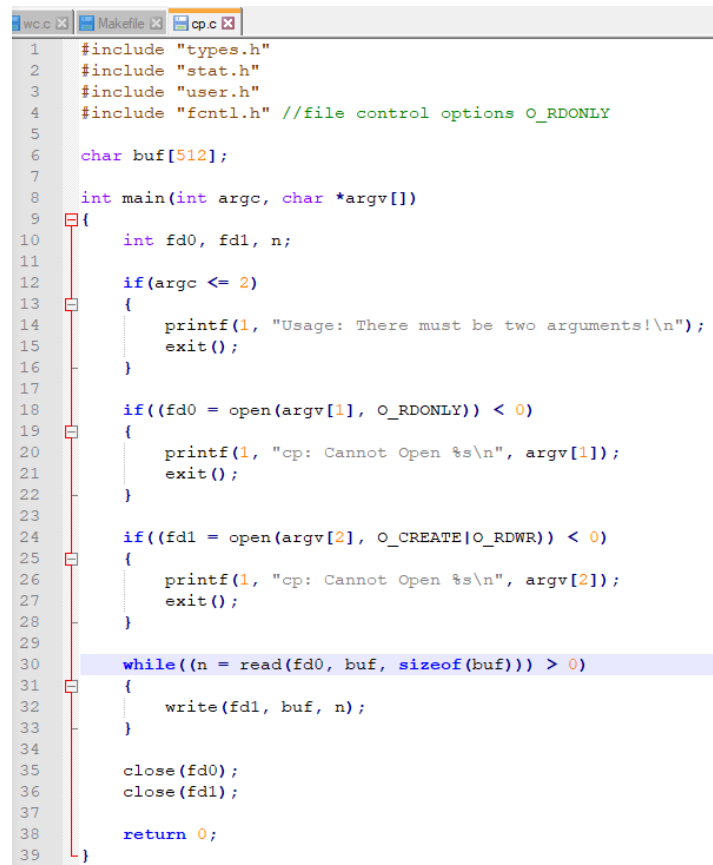
## User Command

### Copy Command

The copy command allows the user to make a duplicate of a file.

#### *cp.c*

The cp command executes the cp.c code to copy a source file as a destination file. The copy command uses command line arguments. The arguments that are used are the source file to be copied and what the destination file should be. If there are not enough arguments, an error message will be displayed. Otherwise, the first file will have its characters read and written to the destination file one by one.



```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4 #include "fcntl.h" //file control options O_RDONLY
5
6 char buf[512];
7
8 int main(int argc, char *argv[])
9 {
10     int fd0, fd1, n;
11
12     if(argc <= 2)
13     {
14         printf(1, "Usage: There must be two arguments!\n");
15         exit();
16     }
17
18     if((fd0 = open(argv[1], O_RDONLY)) < 0)
19     {
20         printf(1, "cp: Cannot Open %s\n", argv[1]);
21         exit();
22     }
23
24     if((fd1 = open(argv[2], O_CREATE|O_RDWR)) < 0)
25     {
26         printf(1, "cp: Cannot Open %s\n", argv[2]);
27         exit();
28     }
29
30     while((n = read(fd0, buf, sizeof(buf))) > 0)
31     {
32         write(fd1, buf, n);
33     }
34
35     close(fd0);
36     close(fd1);
37
38     return 0;
39 }
```

### Makefile

To add the new command to the XV6, the Makefile must be modified to compile the new file. In the Makefile, the user command must be added to “UPROGS” and the file name must be added to “EXTRA”. After this is done, the OS will compile with the new capability for the user.

```
252 EXTRA=\
253     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
254     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c hello.c cp.c\
255     printf.c umalloc.c\
256     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
257     .gdbinit.tmpl gdbutil\
```

```
168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _usertests\
182     _wc\
183     _zombie\
184     _hello\
185     _cp\
```

## System Call

### Current Processes Status

The Current Processes Status (CPS) will show the status of all the current processes. The processes will typically be in three different states of RUNNING, RUNNABLE, or SLEEPING.

*ps.c*

The ps.c file contains the user command that will result in the system call to cps.

```
C: > Users > pakum > AppData > Local > Temp > scp0
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4  #include "fcntl.h"
5
6  int main(int args, char *argv[])
7  {
8      cps();
9      exit();
10 }
```

*syscall.h*

This file contains the declarations for the system call to cps and gives it the number 23.

```
C: > Users > pakum > AppData > Local >
1  // System call numbers
2  #define SYS_fork 1
3  #define SYS_exit 2
4  #define SYS_wait 3
5  #define SYS_pipe 4
6  #define SYS_read 5
7  #define SYS_kill 6
8  #define SYS_exec 7
9  #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_hello 22
24 #define SYS_cps 22
```

This has been changed to 23  
instead of being 22

*syscall.c*

This file contains the available system calls, and states that the `sys_cps` system call will be defined externally.

```
85  extern int sys_chdir(void);
86  extern int sys_close(void);
87  extern int sys_dup(void);
88  extern int sys_exec(void);
89  extern int sys_exit(void);
90  extern int sys_fork(void);
91  extern int sys_fstat(void);
92  extern int sys_getpid(void);
93  extern int sys_kill(void);
94  extern int sys_link(void);
95  extern int sys_mkdir(void);
96  extern int sys_mknod(void);
97  extern int sys_open(void);
98  extern int sys_pipe(void);
99  extern int sys_read(void);
100 extern int sys_sbrk(void);
101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106 extern int sys_hello(void);
107 extern int sys_cps(void);
108
109
110 static int (*syscalls[])(void) =
111 {
112     [SYS_fork]    sys_fork,
113     [SYS_exit]    sys_exit,
114     [SYS_wait]    sys_wait,
115     [SYS_pipe]    sys_pipe,
116     [SYS_read]    sys_read,
117     [SYS_kill]    sys_kill,
118     [SYS_exec]    sys_exec,
119     [SYS_fstat]   sys_fstat,
120     [SYS_chdir]   sys_chdir,
121     [SYS_dup]     sys_dup,
122     [SYS_getpid]  sys_getpid,
123     [SYS_sbrk]    sys_sbrk,
124     [SYS_sleep]   sys_sleep,
125     [SYS_uptime]  sys_uptime,
126     [SYS_open]    sys_open,
127     [SYS_write]   sys_write,
128     [SYS_mknod]   sys_mknod,
129     [SYS_unlink]  sys_unlink,
130     [SYS_link]    sys_link,
131     [SYS_mkdir]   sys_mkdir,
132     [SYS_close]   sys_close,
133     [SYS_hello]   sys_hello,
134     [SYS_cps]     sys_cps,
135 };
```

*defs.h*

The cps function is declared, but not yet defined. The system call sys\_cps will utilize this cps

```

104 //PAGEBREAK: 16
105 // proc.c
106 int      cpuid(void);
107 void     exit(void);
108 int      fork(void);
109 int      growproc(int);
110 int      kill(int);
111 struct cpu* mycpu(void);
112 struct proc* myproc();
113 void     pinit(void);
114 void     procdump(void);
115 void     scheduler(void) __attribute__((noreturn));
116 void     sched(void);
117 void     setproc(struct proc*);
118 void     sleep(void*, struct spinlock*);
119 void     userinit(void);
120 int      wait(void);
121 void     wakeup(void*);
122 void     yield(void);
123 int      cps(void);

```

*user.h*

The cps function is also declared to be callable by the user.

```

C: > Users > pakum > AppData > Local > Temp > scp43770 > hor
1  struct stat;
2  struct rtcdate;
3
4  // system calls
5  int fork(void);
6  int exit(void) __attribute__((noreturn));
7  int wait(void);
8  int pipe(int*);
9  int write(int, const void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int);
13 int exec(char*, char**);
14 int open(const char*, int);
15 int mknod(const char*, short, short);
16 int unlink(const char*);
17 int fstat(int fd, struct stat*);
18 int link(const char*, const char*);
19 int mkdir(const char*);
20 int chdir(const char*);
21 int dup(int);
22 int getpid(void);
23 char* sbrk(int);
24 int sleep(int);
25 int uptime(void);
26 int hello(void);
27 int cps(void);

```

*proc.c*

The actual code to show programs status is defined here as a process. It will allow interrupts so other processes can run while this process is executing. To prevent any of the data from being modified on the process table as it is being read; A lock is added on the ptable and all the process in the table are iterated over, and their statuses are printed. After they are iterated over, the process table is unlocked, allowing modification.

```
535 // Current Process Status
536 int cps()
537 {
538     struct proc *p;
539
540     // Enable interrupts on this processor
541     sti();
542
543     // Loop over process table looking for process to run.
544     acquire(&ptable.lock);
545
546     // Critical Region
547     cprintf("name \t pid \t state \t \t");
548
549     // Print the status of each process
550     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
551     {
552         if(p->state == SLEEPING)
553         {
554             cprintf("%s \t %d \t SLEEPING \t \t", p->name, p->id);
555         }
556         else if(p->state == RUNNING)
557         {
558             cprintf("%s \t %d \t RUNNING \t \t", p->name, p->id);
559         }
560         else if(p->state == RUNNABLE)
561         {
562             cprintf("%s \t %d \t RUNNABLE \t \t", p->name, p->id);
563         }
564     }
565
566     release(&ptable.lock);
567
568     return 23;
569 }
```

*sysproc.c*

This file will define the function `sys_cps`, which is used as the system call. The function simply calls the `cps` function that is defined in the `proc.c`. The screenshot has been modified to what is in the textbox.

```

100 //cps
101 int cps(void)
102 {
103     return cps();
104 }

```

```

int sys_cps(void)
{
    return cps();
}

```

*usys.S*

Since the `cps` is a system call, the process must be managed at the kernel level. This is why the `SYSCALL(cps)` is added.

```

C: > Users > pakum > AppData > Local > Temp > scp
1  #include "syscall.h"
2  #include "traps.h"
3
4  #define SYSCALL(name) \
5      .globl name; \
6      name: \
7          movl $SYS_ ## name, %eax; \
8          int $T_SYSCALL; \
9          ret
10
11  SYSCALL(fork)
12  SYSCALL(exit)
13  SYSCALL(wait)
14  SYSCALL(pipe)
15  SYSCALL(read)
16  SYSCALL(write)
17  SYSCALL(close)
18  SYSCALL(kill)
19  SYSCALL(exec)
20  SYSCALL(open)
21  SYSCALL(mknod)
22  SYSCALL(unlink)
23  SYSCALL(fstat)
24  SYSCALL(link)
25  SYSCALL(mkdir)
26  SYSCALL(chdir)
27  SYSCALL(dup)
28  SYSCALL(getpid)
29  SYSCALL(sbrk)
30  SYSCALL(sleep)
31  SYSCALL(uptime)
32  SYSCALL(hello)
33  SYSCALL(cps)

```



*Makefile*

This file is used for compiling the OS. The new C files have to be added in order for them to compile correctly. The ps user command must be added to UPROGS and EXTRA.

```
168  UPROGS=\
169      _cat\
170      _echo\
171      _forktest\
172      _grep\
173      _init\
174      _kill\
175      _ln\
176      _ls\
177      _mkdir\
178      _rm\
179      _sh\
180      _stressfs\
181      _usertests\
182      _wc\
183      _zombie\
184      _hello\
185      _cp\
186      _ps\
```

```
252  EXTRA=\
253      mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
254      ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c hello.c cp.c ps.c\
255      printf.c umalloc.c\
256      README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
257      .gdbinit.tmpl gdbutil\
```

## Execution

### Commands

#### *cp*

The command `cp` is strictly a user command, and the kernel has no knowledge of its existence.

```
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ ls
.          1 1 512
..         1 1 512
README    2 2 2327
cat       2 3 13688
echo      2 4 12692
forktest  2 5 8128
grep      2 6 15564
init      2 7 13280
kill      2 8 12748
ln        2 9 12648
ls        2 10 14832
mkdir     2 11 12828
rm        2 12 12808
sh        2 13 23296
stressfs  2 14 13476
usertests 2 15 56408
wc        2 16 14224
zombie    2 17 12472
hello     2 18 12420
cp        2 19 13296
console   3 20 0
$ cp README README.txt
pid 4 cp: trap l4 err 5 on cpu 1 eip 0xffffffff addr 0xffffffff--kill proc
$ ls
.          1 1 512
..         1 1 512
README    2 2 2327
cat       2 3 13688
echo      2 4 12692
forktest  2 5 8128
grep      2 6 15564
init      2 7 13280
kill      2 8 12748
ln        2 9 12648
ls        2 10 14832
mkdir     2 11 12828
rm        2 12 12808
sh        2 13 23296
stressfs  2 14 13476
usertests 2 15 56408
wc        2 16 14224
zombie    2 17 12472
hello     2 18 12420
cp        2 19 13296
console   3 20 0
README.txt 2 21 2327
```

#### *ps*

Despite being a user called command. The `ps` command calls the system command of `sys_cps` that eventually will call the `cps` function in the `proc.c`

```
init: starting sh
$ ps
name      pid      state
init      1         SLEEPING
sh        2         SLEEPING
ps        3         RUNNING
$ QEMU: Terminated
```

**Conclusion**

This lab helped show me the difference between a user called process and a system called process. With system calls, many of the system files must be modified to allow the kernel to know about the command and how to handle it. The user commands are simple and easy to implement. Also, I learned more about what how the system files interact with each other. I inadvertently learned more about the zombie command too, which I initially thought was like hello world, but it is entirely different and is used as a demonstration of a process. I also gained some experience in debugging the files that had typos that were hard to detect because they were very close to what they should have been.