

Paul Kummer

Dr. Hanku Lee

Operating Systems CSIS 430-01

10/27/2021

## Implementing process priority

### Purpose

The purpose of this lab is to give the students a way of assigning, changing, and viewing a process's priority. Also, students will reinforce some of the knowledge they have already learned by implementing new user and system calls. Furthermore, the system call that is to be implemented requires the use of a command line argument, which was learned in the previous labs.

### Code

ps command

*proc.h*

```

37 // Per-process state
38 struct proc {
39     uint sz;                // Size of process memory (bytes)
40     pde_t* pgdir;          // Page table
41     char *kstack;          // Bottom of kernel stack for this process
42     enum procstate state;   // Process state
43     int pid;               // Process ID
44     struct proc *parent;    // Parent process
45     struct trapframe *tf;   // Trap frame for current syscall
46     struct context *context; // swtch() here to run process
47     void *chan;            // If non-zero, sleeping on chan
48     int killed;            // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;      // Current directory
51     char name[16];          // Process name (debugging)
52     int priority;           // Process priority (0-50) higher num, higher priority
53 };

```

This file has been modified so that processes now have a priority level.

*proc.c*

```

545 // Current Process Status
546 int cps()
547 {
548     struct proc *p;
549
550     // Enable interrupts on this processor
551     sti();
552
553     // Loop over process table looking for process to run.
554     acquire(&ptable.lock);
555
556     // Critical Region
557     cprintf("name \t pid \t state \t\t priority \t \n");
558
559     // Print the status of each process
560     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
561     {
562         if(p->state == SLEEPING)
563         {
564             cprintf("%s \t %d \t SLEEPING \t %d \n", p->name, p->pid, p->priority);
565         }
566         else if(p->state == RUNNING)
567         {
568             cprintf("%s \t %d \t RUNNING \t %d \n", p->name, p->pid, p->priority);
569         }
570         else if(p->state == RUNNABLE)
571         {
572             cprintf("%s \t %d \t RUNNABLE \t %d \n", p->name, p->pid, p->priority);
573         }
574     }
575
576     release(&ptable.lock);
577
578     return 23;
579 }

```

the cps() function has been modified to now display each processes priority level on execution.

## myfork command

*myfork.c*

```

6  int main(int argc, char* argv[])
7  {
8      int k, n, id;
9      double x = 0, z;
10     double d = 1;
11
12     if(argc < 2)
13     {
14         n = 1;
15     }
16     else
17     {
18         n = atoi(argv[1]);
19     }
20
21     if(n < 0 || n > 20)
22     {
23         n = 2;
24     }
25
26     x = 0;
27     id = 0;
28
29     for(k = 0; k < n; k++)
30     {
31         id = fork();
32
33         if(id < 0)
34         {
35             printf(1, "%d failed in fork! \n", getpid());
36         }
37
38         else if(id == 0)
39         {
40             printf(1, "Child %d created \n", getpid());
41
42             for(z = 0; z < 1000000000.0; z += d)
43             {
44                 x = x + 3.14 * 200.19;
45             }
46
47             break;
48         }
49
50         else
51         {
52             printf(1, "Parent %d creating child %d\n", getpid(), id);
53             wait();
54         }
55     }
56
57     exit();
58 }

```

```

$ myfork && myfork &&
$ Parent Child 10 created
Parent 6 creating child 9
8 creating child 10
Child 9 created
ps
name      pid      state      priority
init      1        SLEEPING   10
sh        2        SLEEPING   10
myfork    9        RUNNING    10
myfork    10       RUNNABLE    10
myfork    6        SLEEPING   10
myfork    8        SLEEPING   10
ps        11       RUNNING    10
$ ps
name      pid      state      priority
init      1        SLEEPING   10
sh        2        SLEEPING   10
myfork    9        RUNNING    10
myfork    10       RUNNABLE    10
myfork    6        SLEEPING   10
myfork    8        SLEEPING   10
ps        12       RUNNING    10
$ ps
name      pid      state      priority
init      1        SLEEPING   10
sh        2        SLEEPING   10
myfork    9        RUNNING    10
myfork    10       RUNNABLE    10
myfork    6        SLEEPING   10
myfork    8        SLEEPING   10
ps        13       RUNNING    10
$ ps
name      pid      state      priority
init      1        SLEEPING   10
sh        2        SLEEPING   10
myfork    9        RUNNING    10
myfork    10       RUNNABLE    10
myfork    6        SLEEPING   10
myfork    8        SLEEPING   10
ps        14       RUNNING    10
$ ps
name      pid      state      priority
init      1        SLEEPING   10
sh        2        SLEEPING   10
myfork    9        RUNNING    10
myfork    10       RUNNABLE    10
myfork    6        SLEEPING   10
myfork    8        SLEEPING   10
ps        15       RUNNING    10

```

myfork will duplicate a parent process n number of times, up to 20. Each child process will be delayed with its execution because it will perform a trivial calculation. The parent must then wait for the child's operation before continuing. This will keep the processes active, allowing for some tests on scheduling and process creation.

cpr command

*cpr.c*

```

C cpr.c > main(int, char *[])
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4  #include "fcntl.h"
5
6  int main(int argc, char* argv[])
7  {
8      int priority, pid;
9
10     if(argc < 3)
11     {
12         printf(2, "Usage: pid priority\n");
13         exit();
14     }
15
16     pid = atoi(argv[1]); //process id to change
17     priority = atoi(argv[2]); //new priority level
18
19     if(priority < 0 || priority > 50)
20     {
21         printf(2, "Invalid priority (0-50)!\n");
22         exit();
23     }
24
25     printf(1, " pid=%d, pr=%d\n", pid, priority);
26     chpr(pid, priority);
27
28     exit();
29 }

```

This file will take in two arguments and change a processes priority. The first argument is the process id to change, and the second argument is the priority to assign to the process. If both values are valid, then chpr will be called to assign the new priority.

*proc.c*

```

581 //Change process priority
582 int chpr(int pid, int priority)
583 {
584     struct proc *p;
585
586     acquire(&ptable.lock);
587
588     for(p=ptable.proc; p<&ptable.proc[NPROC]; p++)
589     {
590         if(p->pid == pid)
591         {
592             p->priority = priority;
593             break;
594         }
595     }
596
597     release(&ptable.lock);
598     return pid;
599 }

```

The function call that acquires a lock on a process and changes its priority. It will return the pid of the current process. This function is used internally by the OS and not directly called by the user.

sysproc.c

```

109 //change process priority
110 int sys_chpr(void)
111 {
112     int pid, pr;
113
114     if(argint(0, &pid) < 0)
115     {
116         return -1;
117     }
118
119     if(argint(1, &pr) < 0)
120     {
121         return -1;
122     }
123
124     return chpr(pid, pr);
125 }

```

The actual system call by the kernel to change a process' priority. It changes the priority if two valid arguments are supplied using the argint() function. Then it calls chpr().

syscall.h

```

1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat    8
10 #define SYS_chdir    9
11 #define SYS_dup    10
12 #define SYS_getpid    11
13 #define SYS_sbrk    12
14 #define SYS_sleep    13
15 #define SYS_uptime    14
16 #define SYS_open    15
17 #define SYS_write    16
18 #define SYS_mknod    17
19 #define SYS_unlink    18
20 #define SYS_link    19
21 #define SYS_mkdir    20
22 #define SYS_close    21
23 #define SYS_hello    22
24 #define SYS_cps    23
25 #define SYS_chpr    24

```

The SYS\_chpr is added as a usable system call with the call number 24.

*defs.h*

```

104 //PAGEBREAK: 16
105 // proc.c
106 int      cpuid(void);
107 void     exit(void);
108 int      fork(void);
109 int      growproc(int);
110 int      kill(int);
111 struct cpu* mycpu(void);
112 struct proc* myproc();
113 void     pinit(void);
114 void     procdump(void);
115 void     scheduler(void) __attribute__((noreturn));
116 void     sched(void);
117 void     setproc(struct proc*);
118 void     sleep(void*, struct spinlock*);
119 void     userinit(void);
120 int      wait(void);
121 void     wakeup(void*);
122 void     yield(void);
123 int      cps(void);
124 int      hello(char*);
125 int      chpr(int, int);

```

chpr() is declared to be a function that will be defined in proc.c.

*usys.S*

```

100 usys.S
1  #include "syscall.h"
2  #include "traps.h"
3
4  #define SYSCALL(name) \
5      .globl name; \
6      name: \
7          movl $SYS_ ## name, %eax; \
8          int $T_SYSCALL; \
9          ret
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
19 SYSCALL(exec)
20 SYSCALL(open)
21 SYSCALL(mknod)
22 SYSCALL(unlink)
23 SYSCALL(fstat)
24 SYSCALL(link)
25 SYSCALL(mkdir)
26 SYSCALL(chdir)
27 SYSCALL(dup)
28 SYSCALL(getpid)
29 SYSCALL(sbrk)
30 SYSCALL(sleep)
31 SYSCALL(uptime)
32 SYSCALL(hello)
33 SYSCALL(cps)
34 SYSCALL(chpr)

```

chpr is added to the list of systemcalls.

*syscall.c*

```

85  extern int sys_chdir(void);
86  extern int sys_close(void);
87  extern int sys_dup(void);
88  extern int sys_exec(void);
89  extern int sys_exit(void);
90  extern int sys_fork(void);
91  extern int sys_fstat(void);
92  extern int sys_getpid(void);
93  extern int sys_kill(void);
94  extern int sys_link(void);
95  extern int sys_mkdir(void);
96  extern int sys_mknod(void);
97  extern int sys_open(void);
98  extern int sys_pipe(void);
99  extern int sys_read(void);
100 extern int sys_sbrk(void);
101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106 extern int sys_hello(void);
107 extern int sys_cps(void);
108 extern int sys_chpr(void);

111 static int (*syscalls[])(void) =
112 {
113     [SYS_fork]    sys_fork,
114     [SYS_exit]    sys_exit,
115     [SYS_wait]    sys_wait,
116     [SYS_pipe]    sys_pipe,
117     [SYS_read]    sys_read,
118     [SYS_kill]    sys_kill,
119     [SYS_exec]    sys_exec,
120     [SYS_fstat]   sys_fstat,
121     [SYS_chdir]   sys_chdir,
122     [SYS_dup]     sys_dup,
123     [SYS_getpid]  sys_getpid,
124     [SYS_sbrk]    sys_sbrk,
125     [SYS_sleep]   sys_sleep,
126     [SYS_uptime]  sys_uptime,
127     [SYS_open]    sys_open,
128     [SYS_write]   sys_write,
129     [SYS_mknod]   sys_mknod,
130     [SYS_unlink]  sys_unlink,
131     [SYS_link]    sys_link,
132     [SYS_mkdir]   sys_mkdir,
133     [SYS_close]   sys_close,
134     [SYS_hello]   sys_hello,
135     [SYS_cps]     sys_cps,
136     [SYS_chpr]    sys_chpr
137 };

```

The system call that calls chpr is declared to be defined elsewhere and the system call itself is added to the array of system calls.

*Makefile*

```

168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _usertests\
182     _wc\
183     _zombie\
184     _hello\
185     _cp\
186     _ps\
187     _testcase\
188     _myfork\
189     _cpr\

255 EXTRA=\
256     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
257     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c\
258     hello.c cp.c ps.c testcase.c myfork.c cpr.c\
259     printf.c umalloc.c\
260     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
261     .gdbinit.tmpl gdbutil\

```

cpr.c is added to the list of files that need to be compiled and linked.

## Code Execution

```

init: starting sh
$ myfork &; myfork &;
$ Parent 7 creatChild 8 created
ng cParent 5 creating child 9
hildChild 9 created
8
ps
name      pid      state      priority
init       1      SLEEPING      10
sh         2      SLEEPING      10
myfork     8      RUNNABLE      10
myfork     9      RUNNING       10
myfork     5      SLEEPING      10
myfork     7      SLEEPING      10
ps         10      RUNNING       10
$ cpr 7 50
pid=7, pr=50
$ ps
name      pid      state      priority
init       1      SLEEPING      10
sh         2      SLEEPING      10
myfork     8      RUNNING       10
myfork     9      RUNNABLE      10
myfork     5      SLEEPING      10
myfork     7      SLEEPING      50
ps         12      RUNNING       10
$ cpr 9 2
pid=9, pr=2
$ ps
name      pid      state      priority
init       1      SLEEPING      10
sh         2      SLEEPING      10
myfork     8      RUNNING       10
myfork     9      RUNNABLE       2
myfork     5      SLEEPING      10
myfork     7      SLEEPING      50
ps         14      RUNNING       10

```

In the execution, first myfork is called twice to create some processes that are running and want to run. Then ps is called to show what processes are running and what the process' priority are. Finally, some of the processes have their priorities changed by using their pid and a new priority value. The results of the change are reflected with the use of the ps command again.

## Conclusion

This lab was useful for understanding how the OS handles processes. It also was good to continue working with system calls and user calls. By repeating these tasks, I am more likely to remember how to do them in the future. Again, it was good to use a command line argument with a system call because this will be useful for other system calls. Lastly, this lab was most useful for setting up scheduling algorithms. It didn't directly do anything with the scheduler, but if the scheduler does start to use the priority level of processes, then processes can be deemed more important than others for whatever reason.